

1. Evaluation Metric/Visualization

Weighted F β -score:

A weighted version of F1-score where β determines the importance of recall relative to precision.

$$F_{\beta} = (1 + \beta^2) \frac{\text{Precision} \cdot \text{Recall}}{\beta^2 \cdot \text{Precision} + \text{Recall}}$$

- If recall is more important, use $\beta > 1$ (e.g., $\beta = 2$).
- If precision is more important, use $\beta < 1$ (e.g., $\beta = 0.5$).

Why using Weighted F β -score: Recall is key for retrieval models.

- 4-star (or higher) ratings are labeled as “interested” and 3-star (or lower) as “not interested.” However, some marginal ratings may still be acceptable.
- Ensuring the retrieval model returns enough relevant items makes recall critical.

F1-score:

When it is difficult to determine the best β for Weighted F β -score, standard F1-score can be used by setting thresholds for precision and recall.

Mean Reciprocal Rank (MRR):

- Reciprocal Rank (RR) measures the rank of the first relevant item in the recommended list:

$$RR = \frac{1}{\text{rank of the first relevant item}}$$

- MRR is the average of reciprocal ranks across all users:

$$MRR = \frac{1}{N} \sum_{i=1}^N \frac{1}{\text{rank of the first relevant item for user } i}$$

- Example:
 - User A: Relevant item at rank 3 \rightarrow Reciprocal Rank = $\frac{1}{3}$
 - User B: Relevant item at rank 1 \rightarrow Reciprocal Rank = $\frac{1}{1}$
 - User C: No relevant item \rightarrow Reciprocal Rank = 0
 - $MRR = \frac{1}{3} \left(\frac{1}{3} + 1 + 0 \right) = 0.44$

Visualization:

- **Precision-Recall curves.**
- **Bar chart of MRR for different users.**

2. Advance Metric

a. Coverage

- **Definition:** Measures the proportion of the catalog (items or users) that the system can recommend:

$$\text{Item Coverage} = \frac{\text{Number of unique items recommended}}{\text{Total number of items in the catalog}}$$

- **Importance:** High coverage ensures a variety of items are recommended, not just popular ones.

b. Diversity

- **Definition:** Measures the dissimilarity among items in a single user's recommendation list.
- Use a similarity metric (e.g., cosine similarity) to calculate pairwise distances between items in the list:

$$\text{Diversity} = 1 - \frac{\sum_{i \neq j} \text{Similarity}(i, j)}{\text{Number of item pairs}}$$

- **Why It Matters:** Balances personalization with exploration.

c. Novelty

- **Definition:** Assesses how unexpected or unique recommendations are. Less popular items contribute more to novelty.
- Example Metric: Inverse Popularity Score:

$$\text{Novelty} = \frac{1}{N} \sum_{i=1}^N -\log_2(\text{popularity of item } i)$$

d. Serendipity

- **Definition:** Evaluates how surprising and delightful recommendations are beyond relevance.
- Example: Measure the difference between expected recommendations (e.g., predicted by collaborative filtering) and actual recommendations.

e. User Effort

- Definition: Measures the effort required by users to find relevant items.
- Example Metric: Average position of relevant items in the recommended list.

$$\text{Effort} = \frac{\sum_{i=1}^N \text{Position of relevant item}}{N}$$

f. Fairness

- Definition: Evaluates whether recommendations are equitably distributed among different user groups or item categories.

3. Optimizing Models

a. Improving Item Similarity Metric

1. **Cosine Similarity:** Measure angular similarity (baseline).
2. **Jaccard Index:** Focus on overlap of interaction sets, effective for binary interaction data.
3. **Adjusted Cosine Similarity:** Account for user biases.
4. **Pearson Correlation Coefficient:** Linear correlation for numerical ratings.
5. **Combine Multiple Metrics:** Weight each based on performance contribution.

b. Features Engineering

1. **Enhanced Weighting Schemes:** Assign higher weights to recent interactions or higher ratings.
2. **Incorporate Additional Parameters:** Add user demographics, business categories, or sentiment analysis of reviews. Consider contextual factors like time of interaction, user location, or session data.
3. **Timeliness and Time Decay:**
 - Exclude older reviews or apply decay factors (e.g., exponential decay).
 - Example:

$$Weight = e^{-\lambda \cdot Age\ of\ Interaction}$$

- λ controls the rate of decay (e.g., 0.01 for gradual decay).

c. Negative Sampling

- Why: Enhances model training by focusing on harder-to-classify items.
- Dynamically generate negative samples during training/testing to address bias.

d. Loss Functions

- Experiment with ranking-specific loss functions like Bayesian Personalized Ranking (BPR) or Listwise Loss.

e. Hyperparameter Tuning

- Optimize parameters like learning rate or embedding dimensions via grid search or Bayesian optimization.

4. Testing Template

Stage	Retrieval									
Model	Item Collaborative Filtering									
Dataset	Yelp (sampled)									
Train Size	80% of dataset for normal case									
Evaluation Metric	Metric	Rationale								
	Recall	Most important metric in retrieval								
	Precision	Ensures retrieved items are relevant								
	F1-score	Provides a balanced view of recall and precision								
	Mean Reciprocal Rank	Measures ranking of relevant items								
	Weighted Fβ-score	Adjusts the trade-off for recall and precision dynamically								
Baseline Model	Simplest Item Collaborative Filtering Model									
	- Item Similarity: Cosine Similarity									
	- Prediction Function:									
	$\sum_j like(user, item_j) \times sim(item_j, item)$									
	- Number of Neighbours: k = 100									
	Features: Total of 1 feature is used									
	- Rate in Review Table									
	Results Overview (example):									
	<table><tr><td>Precision</td><td></td></tr><tr><td>Recall</td><td></td></tr><tr><td>F1-Score</td><td></td></tr><tr><td>MRR</td><td></td></tr></table>		Precision		Recall		F1-Score		MRR	
	Precision									
Recall										
F1-Score										
MRR										
Graphs: Precision-Recall Curve, Bar Chart of MRR for Different Users										
Evaluation Strategies/Steps	Strategy 1: K-Nearest Neighbors (KNN)									
	Set-up:									
	1. Prepare user and business index tables.									
	2. Split 80/20 for train/test data.									
	3. Exclude test data during KNN computation.									
Evaluation Strategies/Steps	4. Compute KNN results for evaluation.									
	Evaluation:									
	1. Take a pair of user_id and business_id as input.									
	2. Compute the list of item similarity with the user recent interacted items (some may be 0, max case will have N similarities).									
	3. Compute all the interest scores.									
Evaluation Strategies/Steps	4. Compute the evaluation metric.									
	Precision	True Positive								
	User gave a rate of 4 or higher for business, predicted interest score is in top K results.									

		True Negative	User gave a rate of 3 or lower for business, predicted interest score is not in top K results.
	Recall	True Positive	Same as Precision .
		False Positive	User gave a rate of 3 or below for business, predicted interest score are in top K results.
	F1-Score	Usual Formula	
	Mean Reciprocal Rank	RR	The respective rank of True Positive business in the top K results.
	Weighted Fβ-score	β	$B > 1$ as recall is important here.
	Strategy 2: Leave-One-Out Cross-Validation (LOOCV) Set-up: <ol style="list-style-type: none"> 1. Set up the user and business index tables. 2. For each user, leave one interaction out of the training set and treat it as the test case. 3. Exclude this data from training. 4. Compute KNN results for evaluation. Evaluation: Same metrics as Strategy 1, ensuring robust testing for individual interactions.		